

AI ALIGNMENT ISSUE:


SHOULD WE FOLLOW A UTILITARIAN APPROACH?

CENTRAL THESIS

I argue that AI models should be aligned with a people's relative and contemporary environment because of a simple observation of the current differing goals, morals, and ideals of countries.


VISUALIZED

- 10 values appeared universally
- Differences are in emphasis, not 'kinds of'
- Locally, in the US 90% of voters agree on the same fundamental values with little give (~5% disagreeing on fundamentals)



← Schwartz Value Circumplex

Validated across 80+ countries and 60,000 people. The structure holds everywhere.



John Stuart Mill

MY REASONING

- The main scare of the AI alignment issue deals with the idea that humans are already fundamentally conflicted and have countless internal conflicts
- However, the empirical evidence universally (Schwartz) and in the US (Wolf & Hanel) I have showcased suggest otherwise. There seems to be a feasible target to aim for
- Given a utility function, minimize suffering and maximizing well-being (John Stuart Mill), which is largely agreed upon then given disagreements, such as gun control vs gun ownership, these are not conflicts of values but empirical disputes about the methods we can approach
- So, the approach is to align AI with this substantial common ground we share and build the AI to then defer to humans for contested points and methods
- This is precisely where the Utilitarian approach thrives. The AI we design to align with us humans will be the tool to optimize towards outcomes that we, as a collective, strive for rather than design the AI to impose a value system on others
- Aligning AI to the 90% as a baseline includes the vast majority of the population and leaves the remaining 10% of genuinely contested values to a process that includes human oversight
- An AI which will be grounded in practically universal values would face less resistance and more trust than if the model(s) tried to tackle each individual issue independently
- So, maybe we do not need perfect in-line agreement with one another to achieve a meaningful outcome and alignment. We just need to recognize the common ground that is already there

OPPOSITION

- Critics may argue that these values, like keeping people safe, are too abstract and obscure the irreconcilable differences between people
- Given the example of gun control vs gun ownership to ensure safety, the AI must ultimately choose a direction to approach
- However, this only bolsters the Utilitarian position because the AI will not pick a side, the AI will evaluate which method actually does produce more safety. It becomes a data problem, not a values problem
- Another critical analysis might reveal that there is a danger of the tyrannical majority, the 90% overriding the 10% and in history minority positions have ended up being the factually correct position
- This is why the AI defers to human insight for the contested values rather than allowing the AI to optimize with no oversight

SOURCES

- Wolf, Lukas, and Paul Hanel. "Perceived Polarization and Shared Values: Evidence from US Democrats and Republicans." *Social Psychological and Personality Science*, 2024.
- Schwartz, Shalom H. "An Overview of the Schwartz Theory of Basic Values." *Online Readings in Psychology and Culture*, vol. 2, no. 1, 2012.

